



Institut et hôpital neurologiques de Montréal
Montreal Neurological Institute and Hospital



Jewish General Hospital
Lady Davis Institute for Medical Research



CENTRE
LUDMER
NEUROINFORMATIQUE & SANTÉ MENTALE | NEUROINFORMATICS & MENTAL HEALTH

Intro to Brain Segmentation with Keras

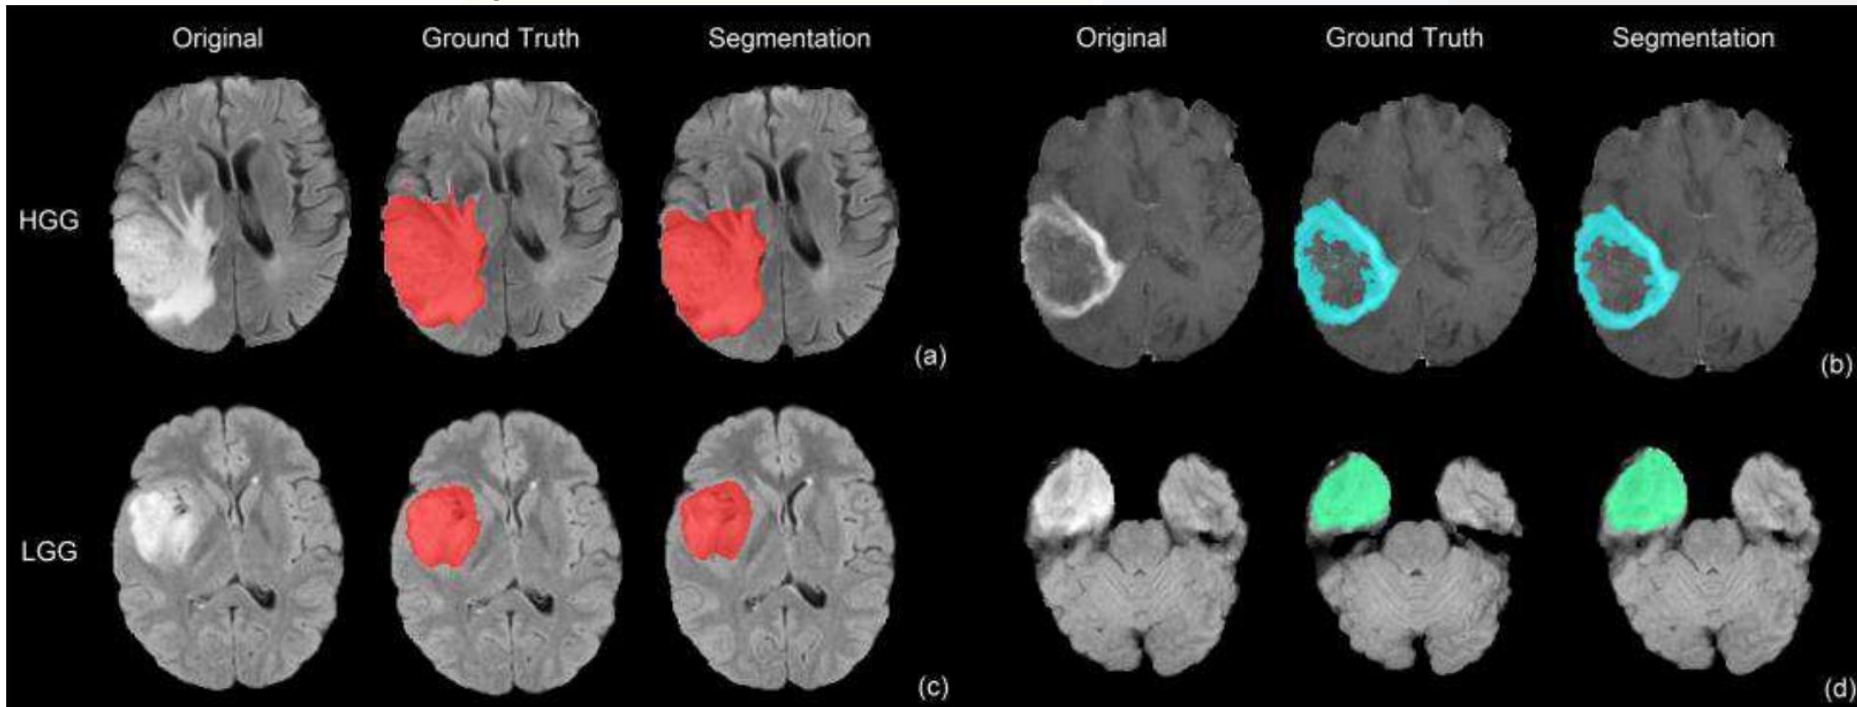
MAIN 2018 Educational Course

Thomas Funck, McGill University

Email: tffunck@gmail.com; Twitter: [@tffunck](https://twitter.com/tffunck)

Why CNNs for segmentation?


- Can solve complicated problems
 - ...if you have the data + labels



Dong, et al. (2017) Automatic Brain Tumor Detection and Segmentation Using U-Net Based Fully Convolutional Networks. <https://arxiv.org/abs/1705.03820>

Why CNNs for segmentation?

- Can solve complicated problems
- Can also solve simple, but annoying problems



Why CNNs for segmentation?

- Can solve complicated problems
- Can also solve simple, but annoying problems
- Goal: Dive into nitty-gritty to building a CNN
 - So that you can go back to the lab and create your own



Outline

1. Basic Concepts
2. Simple Networks
3. U-Net



Outline


1. Basic Concepts

- a. Kernels
- b. Receptive field & dilations
- c. Upsampling & downsampling
- d. Final activation functions
- e. Loss functions
- f. Metrics
- g. Cross-validation
- h. Feature extraction



Example of a simple convolutional network

- 1 Layer, 3 Kernels of 3x3 size




Example of a simple convolutional network

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	(None, 110, 92, 1)	0
conv2d_1 (Conv2D)	(None, 110, 92, 3)	30
dropout_1 (Dropout)	(None, 110, 92, 3)	0
conv2d_2 (Conv2D)	(None, 110, 92, 3)	12

Total params: 42
Trainable params: 42
Non-trainable params: 0

Example of a simple convolutional network

- 3 Layer, 3 Kernels of 3x3 size



Example of a simple convolutional network

- 3 Layer, 3 Kernels of 3x3 size

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	(None, 110, 92, 1)	0
conv2d_3 (Conv2D)	(None, 110, 92, 3)	30
dropout_2 (Dropout)	(None, 110, 92, 3)	0
conv2d_4 (Conv2D)	(None, 110, 92, 3)	84
dropout_3 (Dropout)	(None, 110, 92, 3)	0
conv2d_5 (Conv2D)	(None, 110, 92, 3)	84
dropout_4 (Dropout)	(None, 110, 92, 3)	0
conv2d_6 (Conv2D)	(None, 110, 92, 3)	12

Total params: 210
Trainable params: 210
Non-trainable params: 0





Different number of parameters. What gives?

- 3 Layer, 3 Kernels of 3x3 size

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	(None, 110, 92, 1)	0
conv2d_3 (Conv2D)	(None, 110, 92, 3)	30
dropout_2 (Dropout)	(None, 110, 92, 3)	0
conv2d_4 (Conv2D)	(None, 110, 92, 3)	84
dropout_3 (Dropout)	(None, 110, 92, 3)	0
conv2d_5 (Conv2D)	(None, 110, 92, 3)	84
dropout_4 (Dropout)	(None, 110, 92, 3)	0
conv2d_6 (Conv2D)	(None, 110, 92, 3)	12

Total params: 210
Trainable params: 210
Non-trainable params: 0






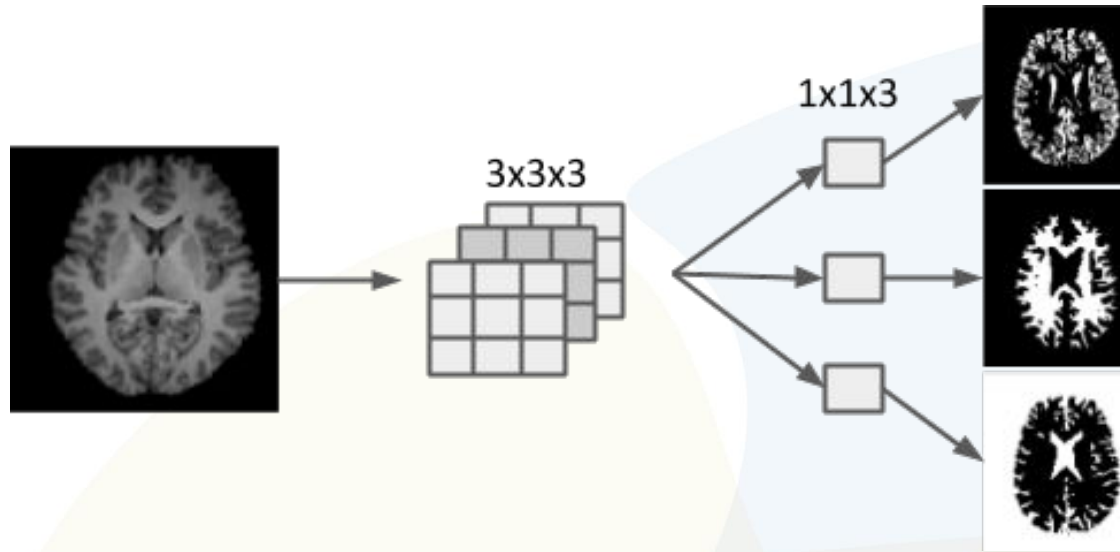
$$N_{parameters} = N_{kernels} \times Kernel_{Dimension1} \times Kernel_{Dimension2} \times Kernel_{Channels} + N_{kernels}$$


$$84 = 3 \times 3 \times 3 \times 3 + 3$$

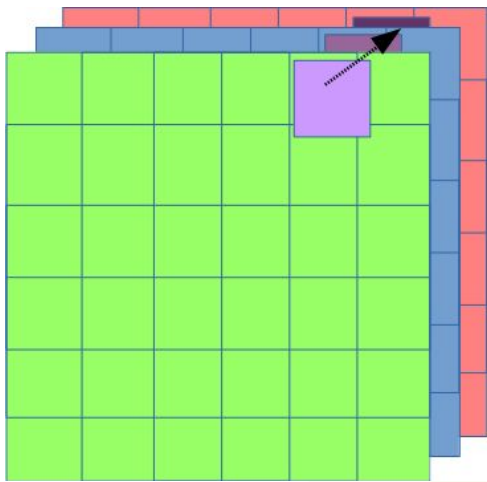
Layer (type)	Output Shape	Param #
input_2 (InputLayer)	(None, 110, 92, 1)	0
conv2d_3 (Conv2D)	(None, 110, 92, 3)	30
dropout_2 (Dropout)	(None, 110, 92, 3)	0
conv2d_4 (Conv2D)	(None, 110, 92, 3)	84
dropout_3 (Dropout)	(None, 110, 92, 3)	0
conv2d_5 (Conv2D)	(None, 110, 92, 3)	84
dropout_4 (Dropout)	(None, 110, 92, 3)	0
conv2d_6 (Conv2D)	(None, 110, 92, 3)	12

Total params: 210
 Trainable params: 210
 Non-trainable params: 0

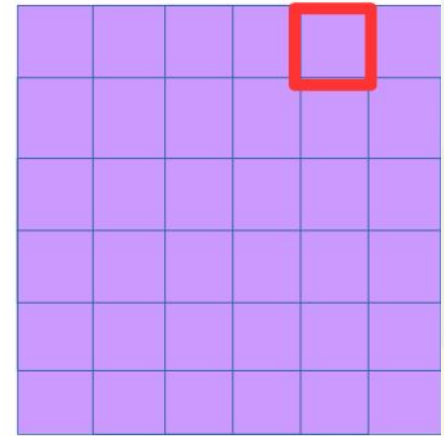







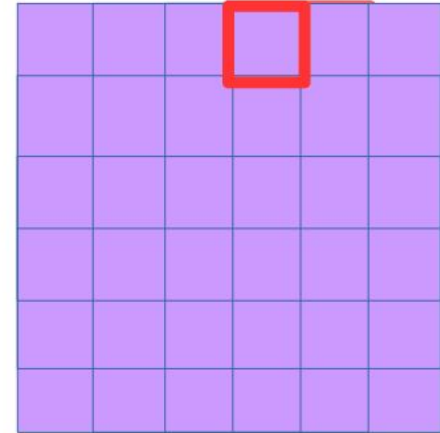


$$\text{Green} \times \text{Purple} + \text{Blue} \times \text{Dark Purple} + \text{Red} \times \text{Dark Purple} =$$



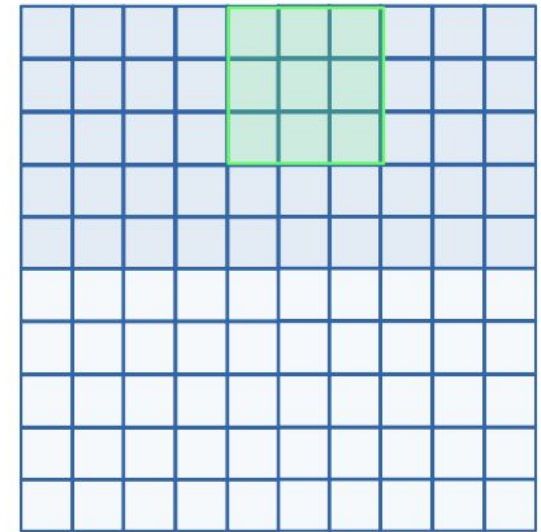




$$\text{Green} \times \text{Purple} + \text{Blue} \times \text{Dark Purple} + \text{Red} \times \text{Darkest Purple} =$$



Kernel Stride


- must divide the image evenly
 - you can pad the image if necessary



kernel stride = 2


Kernel Padding

padding = "valid"




Kernel Padding

padding = "valid"




Kernel Padding

padding = "valid"




Kernel Padding

padding = "same"




Kernel Padding

padding = "same"



Kernel Padding


padding = "same"



Outline

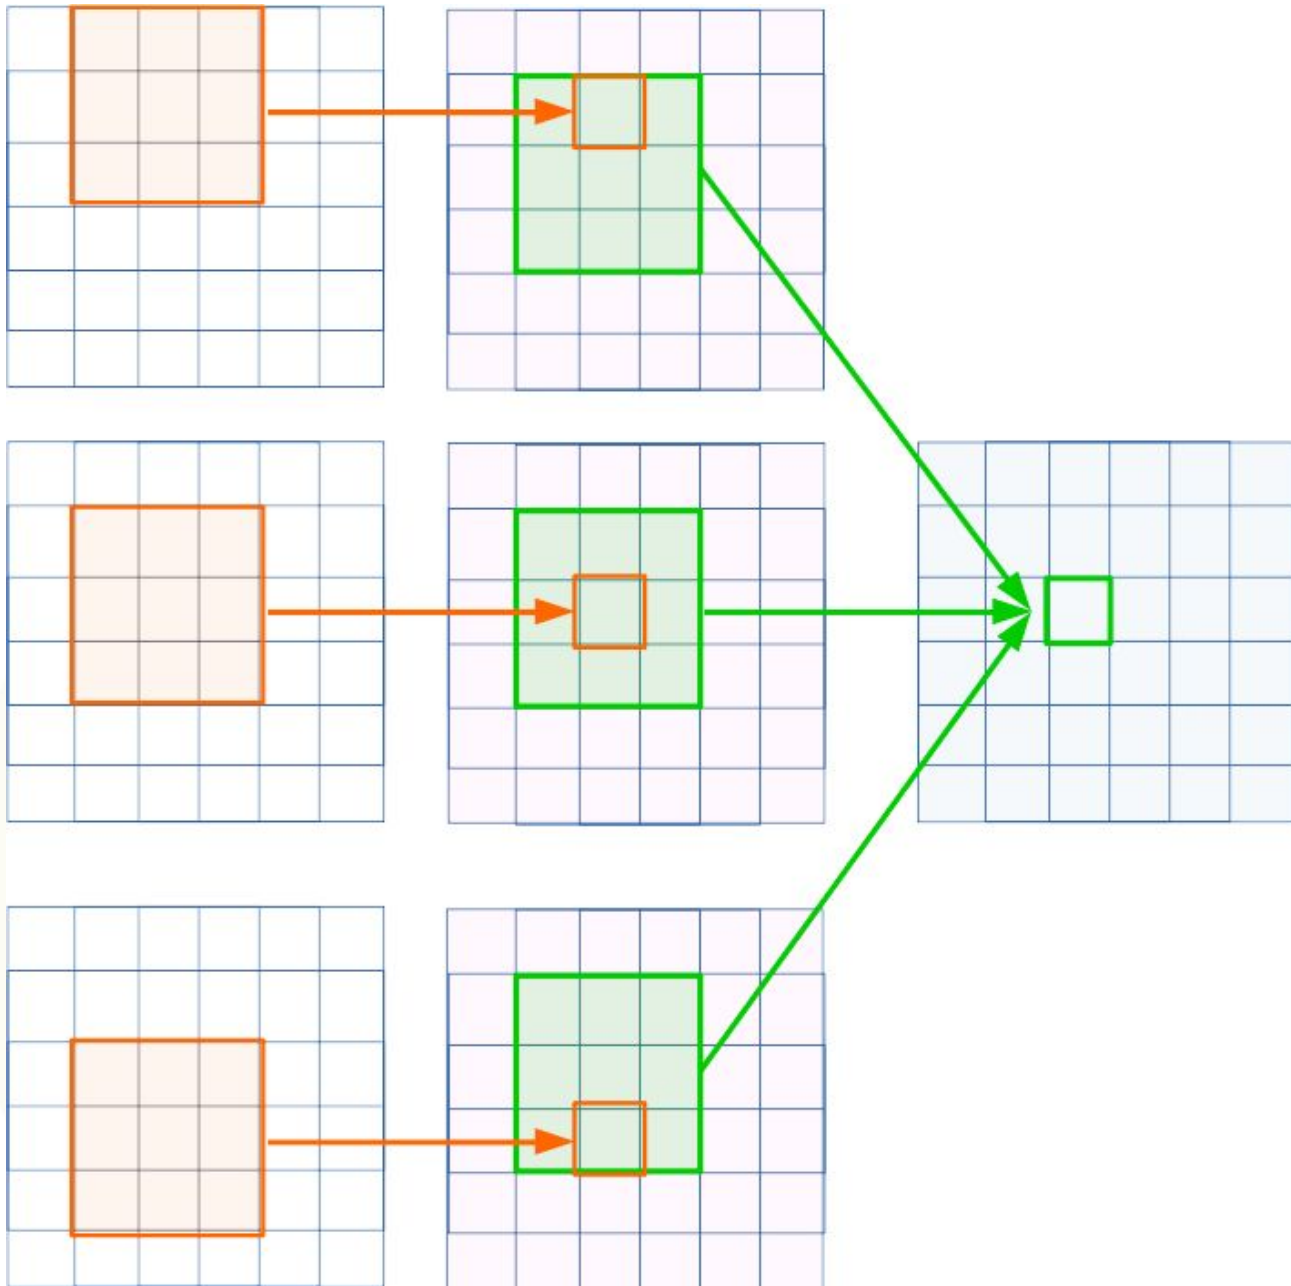
1. Basic Concepts


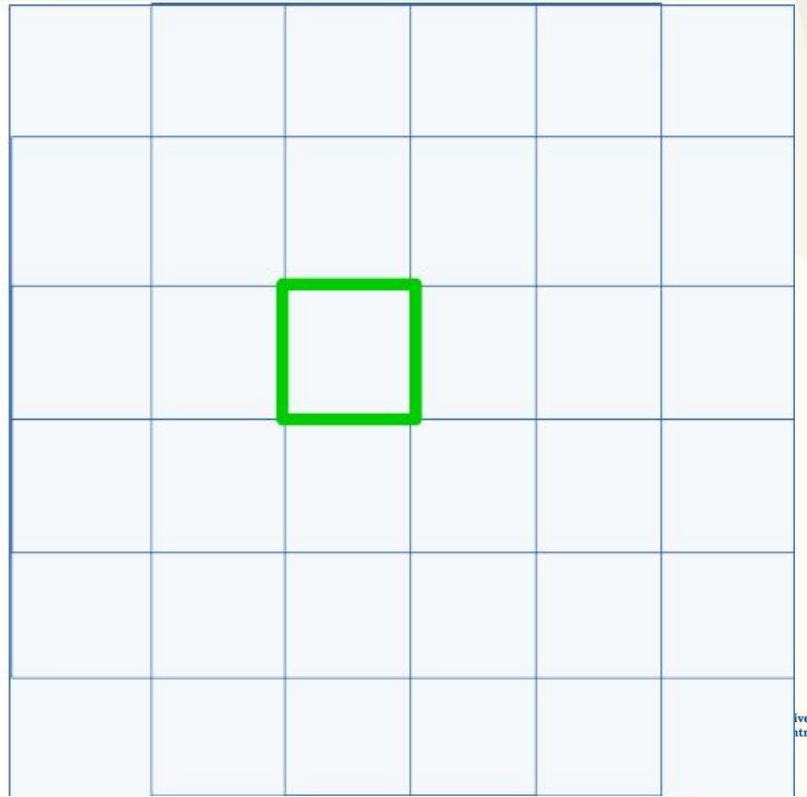
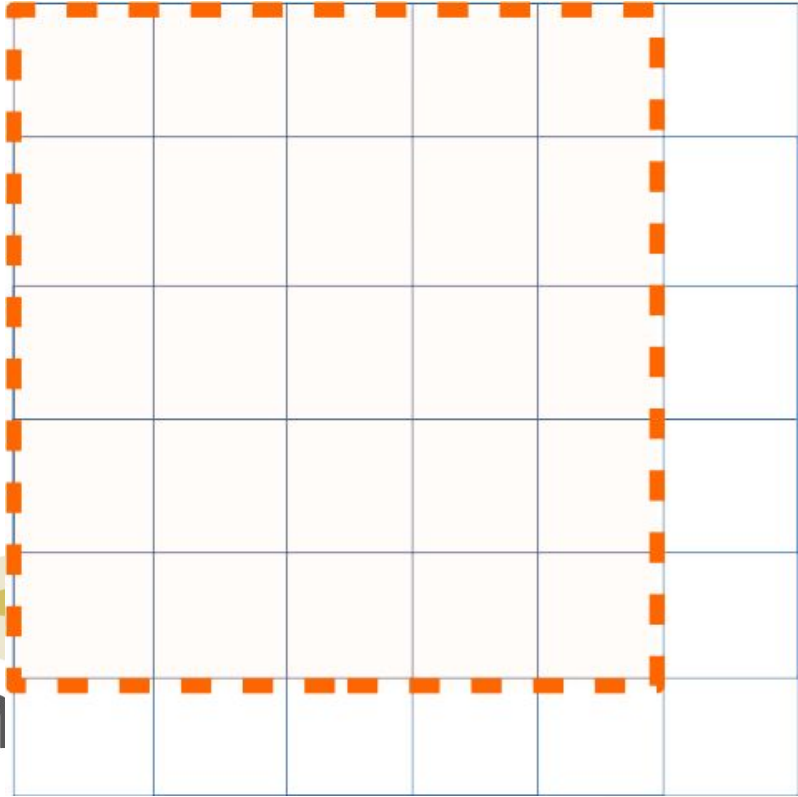
- a. Kernels
- b. Receptive field & dilations
- c. Upsampling & downsampling
- d. Final activation functions
- e. Loss functions
- f. Metrics
- g. Cross-validation
- h. Feature extraction



*“The **receptive field** of an individual sensory neuron is the particular region of the sensory space in which a stimulus will modify the firing of that neuron.”*


https://en.wikipedia.org/wiki/Receptive_field






Dilations

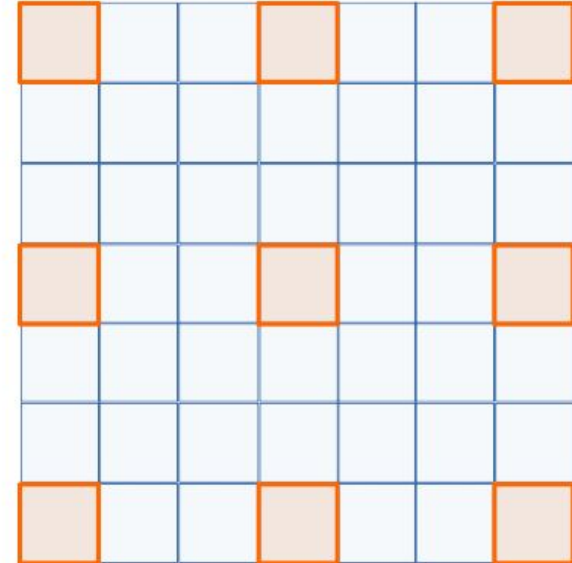
Dilation=1, (3 x 3)



Dilation=2, (3 x 3)



Dilation=3, (3 x 3)



- Increase receptive field without increasing parameters

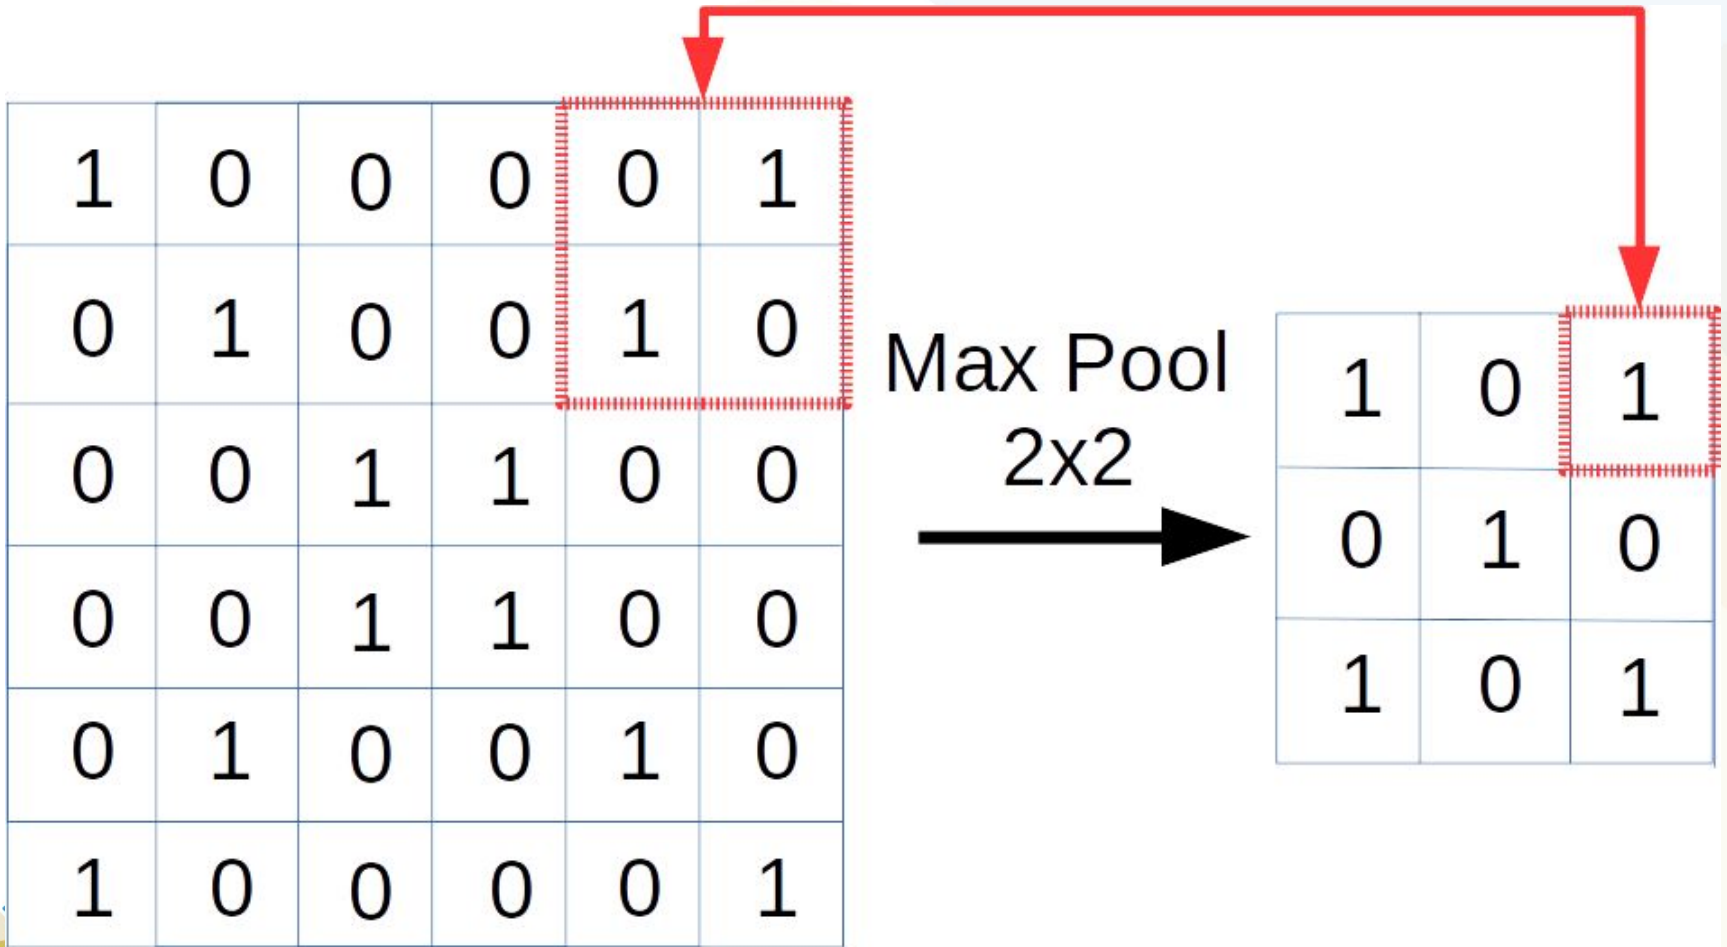
Outline

1. Basic Concepts

- a. Kernels
- b. Receptive field & dilations
- c. Upsampling & downsampling
- d. Final activation functions
- e. Loss functions
- f. Metrics
- g. Cross-validation
- h. Feature extraction

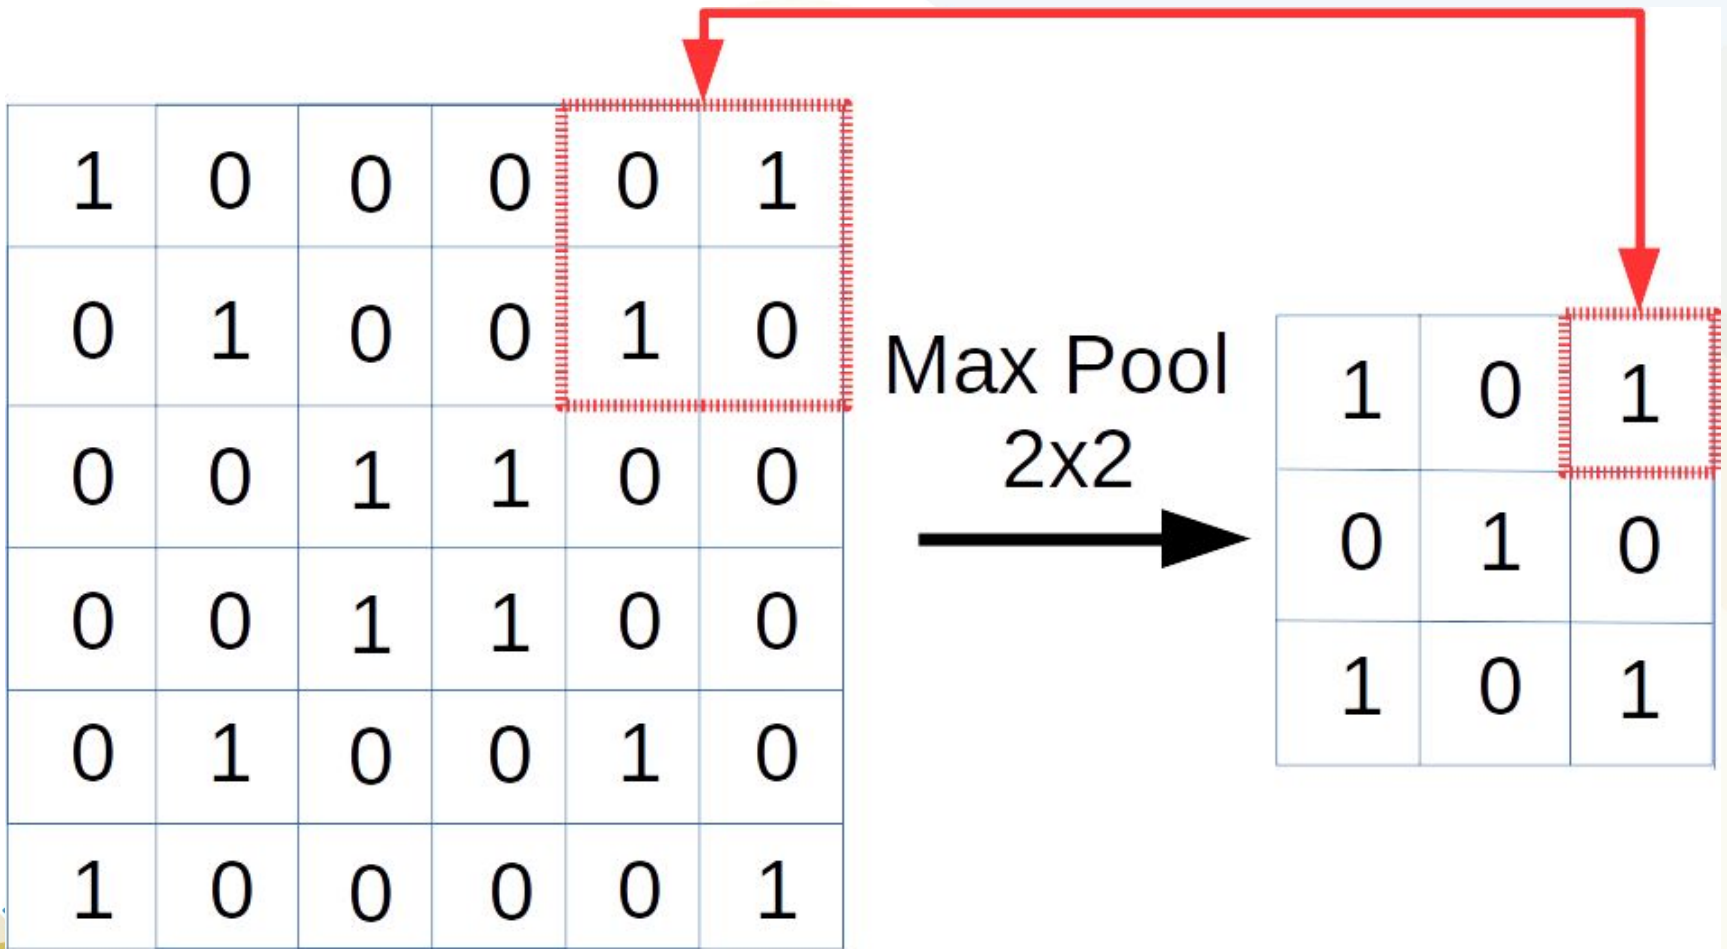


Downsampling



Summarizes essential features of image at lower spatial resolution

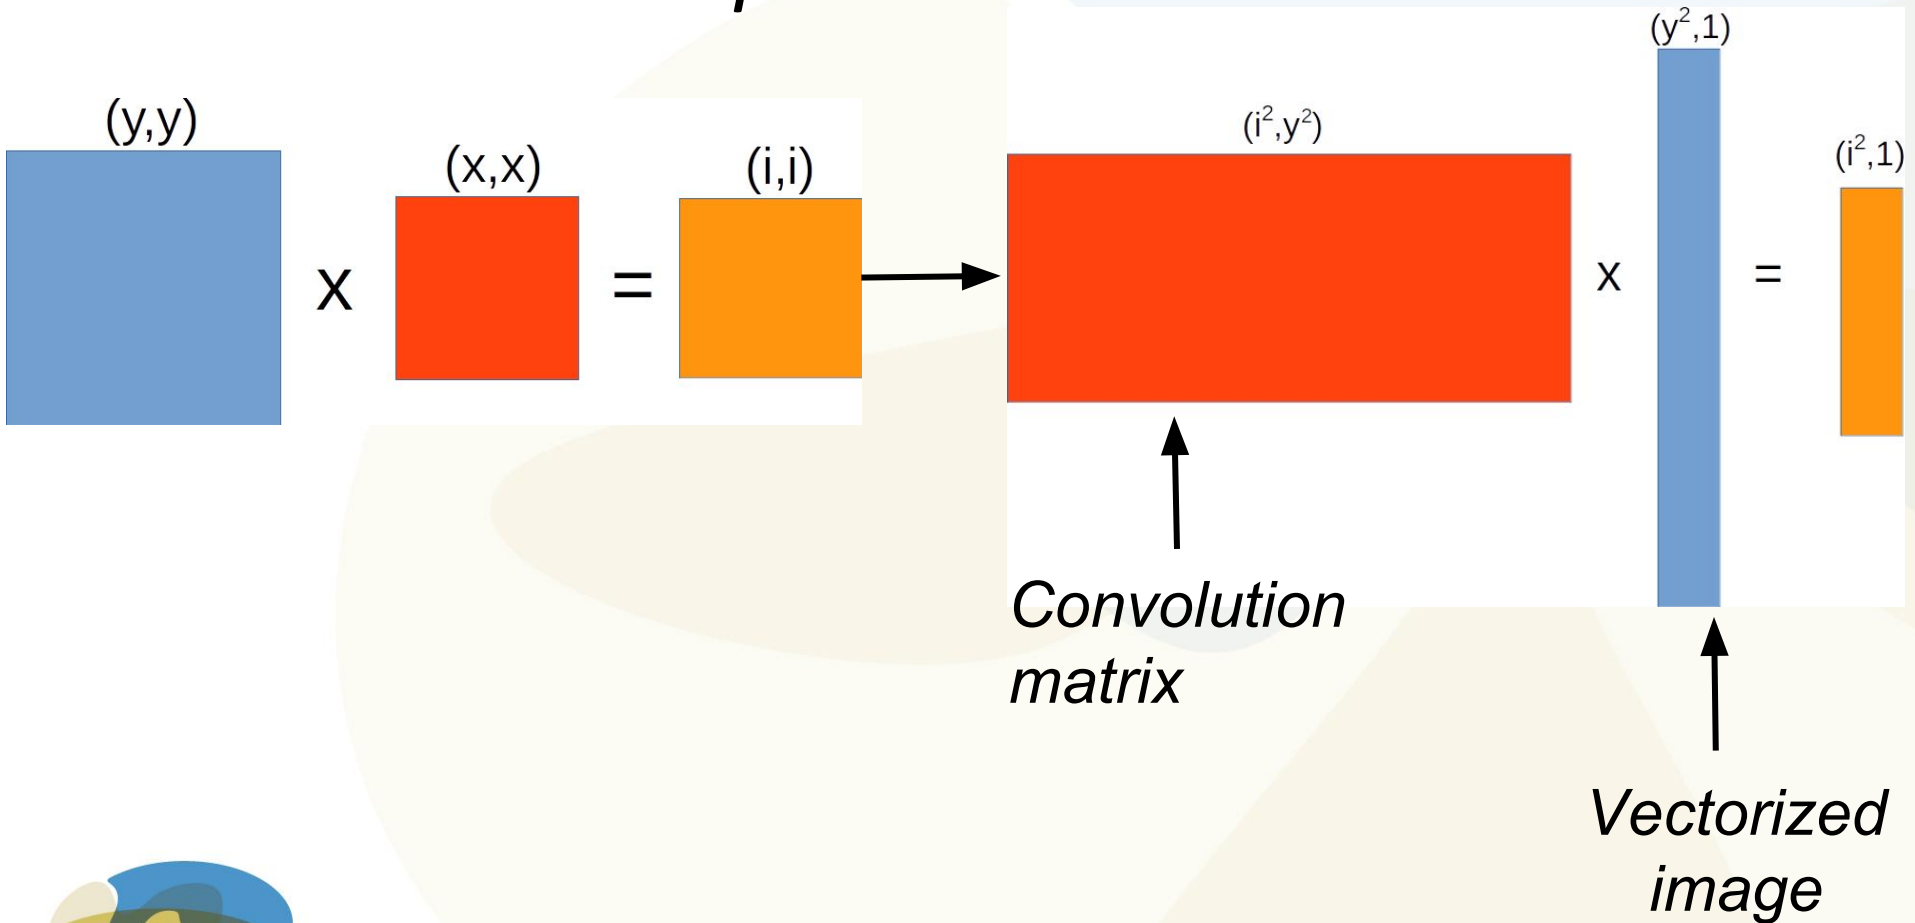
Downsampling



Decreases the image size by half!

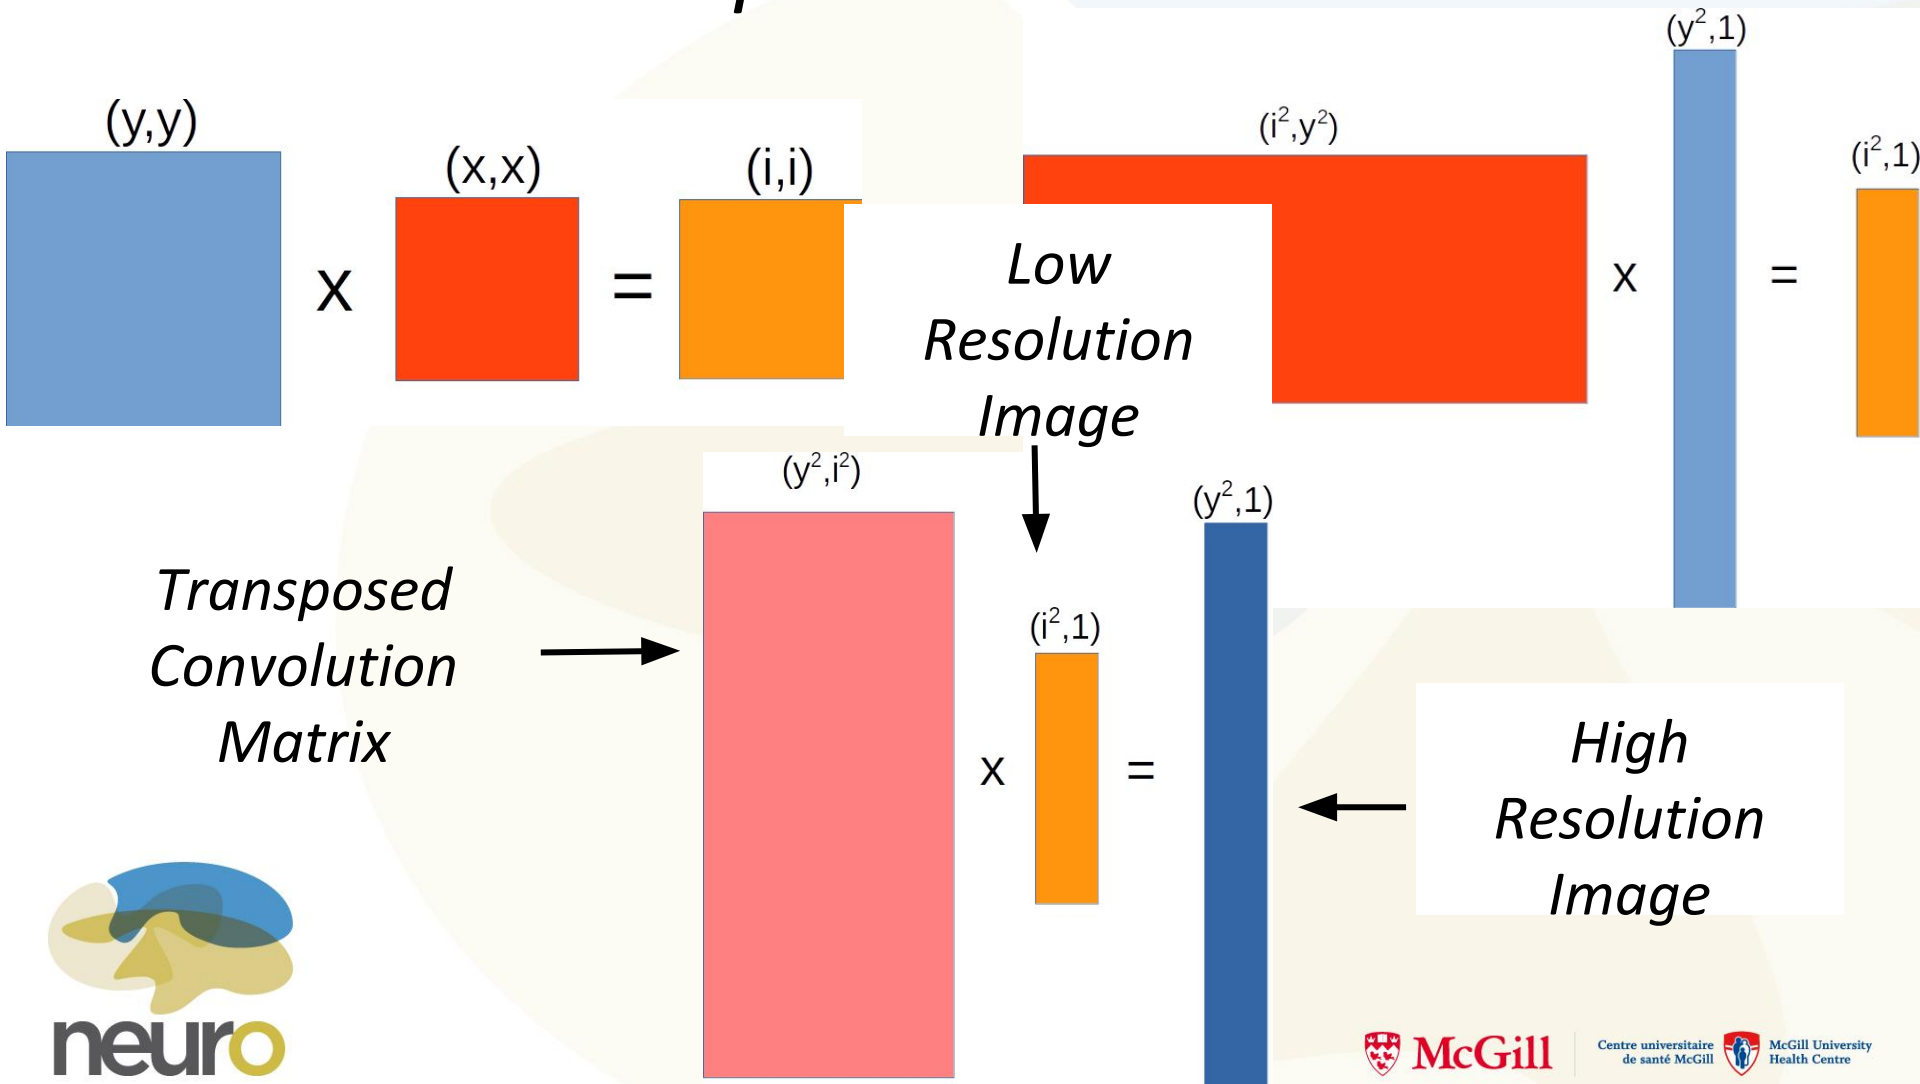
Upsampling

Transpose Convolution



Upsampling

Transpose Convolution




Upsampling

Interpolation + 2x2 Convolution

Interpolation
(nearest neighbours)


2x2 Convolution



Upsampling

Interpolation

Interpolation
(nearest neighbours)




Outline


1. Basic Concepts

- a. Kernels
- b. Receptive field & dilations
- c. Upsampling & downsampling
- d. Final activation functions
- e. Loss functions
- f. Metrics
- g. Cross-validation
- h. Feature extraction






Binary Classification




Binary Classification

Sigmoid :

$$S(x) = \frac{1}{1 + e^{-x}}$$



Multi-category Classification



Multi-category Classification

Softmax :


- Generalization of binary sigmoid
- Creates pseudo-probability distribution

$$\sigma(x_j) = \frac{e^{x_j}}{\sum_i e^{x_i}}$$

Softmax Output of Network

Softmax

- Each pixel = $[P(\text{GM}), P(\text{WM}), P(\text{BG})]$
- 3D output array = (Width, Height, 3)



	GM	WM	CSF
Probability	0.1	0.85	0.05

	GM	WM	CSF
Probability	0.7	0.2	0.1

- Transform to 2D image by finding class with max probability at each pixel

Outline

1. Basic Concepts

- a. Kernels
- b. Receptive field & dilations
- c. Upsampling & downsampling
- d. Final activation functions
- e. Loss functions
- f. Metrics
- g. Cross-validation
- h. Feature extraction



Choosing loss function

- **Cross Entropy :**

- From information theory
- Quantifies difference between two probability distributions

$$CE = \frac{1}{N} \sum_i^N \sum_c^C True_i^c \times \log_2(Predicted_i^c)$$

Choosing loss function

Cross Entropy :

$$CE = \frac{1}{N} \sum_i^N \sum_c^C True_i^c \times \log_2(Predicted_i^c)$$

Ground truth
probability
from labels

Bits of
information
for predicted
label

Cross Entropy

True Distribution

	GM	WM	CSF
<i>data_0</i>	1	0	0
<i>data_1</i>	0	1	0
<i>data_2</i>	0	1	0

Predicted Distribution

	GM	WM	CSF
<i>data_0</i>	.5	.2	.3
<i>data_1</i>	.2	.4	.4
<i>data_2</i>	.1	.5	.4

Predicted distribution produced with softmax / sigmoid final activation

Cross Entropy

True

GM WM CSF

<i>data_0</i>	1	0	0
<i>data_1</i>	0	1	0
<i>data_2</i>	0	1	0

Predicted

GM WM CSF

<i>data_0</i>	.5	.2	.3
<i>data_1</i>	.2	.4	.4
<i>data_2</i>	.1	.5	.4

$$CE = \frac{1}{N} \sum_i^N \sum_c^C True_i^c \times \log_2(Predicted_i^c)$$

$$1 \times -\log_2(0.5) + 0 \times \log_2(.2) + 0 \times \log_2(.3) +$$

Cross Entropy

True

GM WM CSF

<i>data_0</i>	1	0	0
<i>data_1</i>	0	1	0
<i>data_2</i>	0	1	0

Predicted

GM WM CSF

<i>data_0</i>	.5	.2	.3
<i>data_1</i>	.2	.4	.4
<i>data_2</i>	.1	.5	.4

$$CE = \frac{1}{N} \sum_i^N \sum_c^C True_i^c \times \log_2(Predicted_i^c)$$

$$1 \times -\log_2(0.5) + 0 \times \log_2(.2) + 0 \times \log_2(.3) +$$

$$0 \times -\log_2(0.2) + 1 \times \log_2(.4) + 0 \times \log_2(.4) +$$

Cross Entropy

True

GM WM CSF

<i>data_0</i>	1	0	0
<i>data_1</i>	0	1	0
<i>data_2</i>	0	1	0

Predicted

GM WM CSF

<i>data_0</i>	.5	.2	.3
<i>data_1</i>	.2	.4	.4
<i>data_2</i>	.1	.5	.4

$$CE = \frac{1}{N} \sum_i^N \sum_c^C True_i^c \times \log_2(Predicted_i^c)$$

$$1 \times -\log_2(0.5) + 0 \times \log_2(.2) + 0 \times \log_2(.3) +$$

$$0 \times -\log_2(0.2) + 1 \times \log_2(.4) + 0 \times \log_2(.4) +$$

$$0 \times -\log_2(0.1) + 1 \times \log_2(.5) + 0 \times \log_2(.4)$$

Cross Entropy

True

GM WM CSF

<i>data_0</i>	1	0	0
<i>data_1</i>	0	1	0
<i>data_2</i>	0	1	0

Predicted

GM WM CSF

<i>data_0</i>	.5	.2	.3
<i>data_1</i>	.2	.4	.4
<i>data_2</i>	.1	.5	.4

$$CE = -\frac{1}{N} \sum_i^N \sum_c^C True_i^c \times \log_2(Predicted_i^c)$$

$$1 \times -\log_2(0.5) + 0 \times \log_2(.2) + 0 \times \log_2(.3) +$$

$$0 \times -\log_2(0.2) + 1 \times \log_2(.4) + 0 \times \log_2(.4) +$$

$$0 \times -\log_2(0.1) + 1 \times \log_2(.5) + 0 \times \log_2(.4)$$

$$= 1 \times -\log_2(0.5) + 1 \times \log_2(.4) + 1 \times \log_2(.4)$$

$$= 3.64$$

Cross Entropy

True

GM WM CSF

<i>data_0</i>	1	0	0
<i>data_1</i>	0	1	0
<i>data_2</i>	0	1	0

Predicted

GM WM CSF

<i>data_0</i>	.8	.1	.1
<i>data_1</i>	.1	.8	.1
<i>data_2</i>	.1	.8	.1

Cross Entropy

True

GM WM CSF

<i>data_0</i>	1	0	0
<i>data_1</i>	0	1	0
<i>data_2</i>	0	1	0

Predicted

GM WM CSF

<i>data_0</i>	.8	.1	.1
<i>data_1</i>	.1	.8	.1
<i>data_2</i>	.1	.8	.1

$$CE = -\frac{1}{N} \sum_i^N \sum_c^C True_i^c \times \log_2(Predicted_i^c)$$

$$= 1 \times -\log_2(0.8) + 1 \times \log_2(.8) + 1 \times \log_2(.8)$$

$$= 0.97$$

By improving our prediction, we've decreased the value of the cross entropy function

Outline

1. Basic Concepts

- a. Kernels
- b. Receptive field & dilations
- c. Upsampling & downsampling
- d. Final activation functions
- e. Loss functions
- f. Metrics**
- g. Cross-validation
- h. Feature extraction

Metrics

1. Classification labels are integers
2. Dice Metric
 - a. Binary classification
 - b. Custom metric, not included by default in Keras
 - c. Perfect overlap = $(2 \times |1|) / (1+1) = 1$

$$\frac{2|X \cap Y|}{|X| + |Y|}$$

Metrics

1. Classification labels are integers
 - a. → Can't use MSE or similar metrics
2. Dice Metric
3. Categorical Accuracy
 - a. Binary and multi-class labels
 - b. Use 'acc' in Keras

$$\frac{1}{N} \sum_i I(\text{Predicted}_i = \text{Label}_i)$$

Metrics

1. Classification labels are integers
2. Dice Metric
3. Categorical Accuracy
4. Baseline for metrics is not usually not 0
 - a. Ex. Guess all 0 \rightarrow metric = 0.6


Outline

1. Basic Concepts

- a. Kernels
- b. Receptive field & dilations
- c. Upsampling & downsampling
- d. Final activation functions
- e. Loss functions
- f. Metrics
- g. **Cross-validation**
- h. Feature extraction




K-folds cross validation




Data

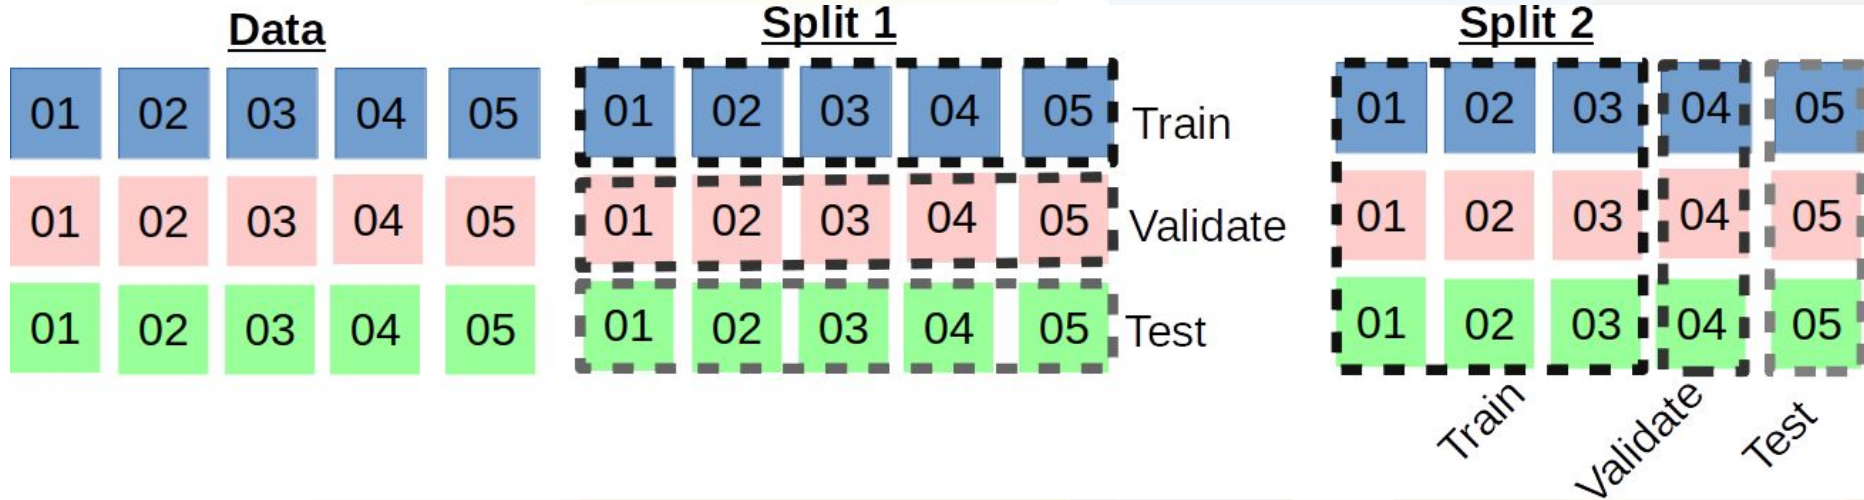
01	02	03	04
05	06	07	08
09	10	11	12
13	14	15	16
17	18	19	20
21	22	23	24

Split 1



Split 2






Outline

1. Basic Concepts

- a. Kernels
- b. Receptive field & dilations
- c. Upsampling & downsampling
- d. Final activation functions
- e. Loss functions
- f. Metrics
- g. Cross-validation
- h. Feature extraction


Many ways to analyze volumes

Full 3D Volume




Many ways to analyze volumes

Full 3D Volume




2D Slices




Many ways to analyze volumes


Full 3D Volume



2D Slices




3D Patches





Formatting data to feed into Keras






3D Nifti / MINC



2D Slices



HDF5 / NPY



More info

1. Overview
 - a. **“Convolutional Neural Nets for Visual Recognition”**, <http://cs231n.github.io/convolutional-networks/>, karpathy@cs.stanford.edu
 - i. *This is a very good overview. Definitely read!*
 - b. **“Deep Learning”** <http://neuralnetworksanddeeplearning.com/chap6.html> , Michael Nielsen
2. Kernels
 - a. **“Convolutional Neural Networks (CNN, or ConvNets)”**, <https://medium.com/@phidaouss/convolutional-neural-networks-cnn-or-convnets-d7c688b0a207>, Firdaouss Doukkali.
3. Softmax
 - a. **“Difference Between Softmax Function and Sigmoid Function”**, <http://dataaspirant.com/2017/03/07/difference-between-softmax-function-and-sigmoid-function/> , Saimadhu Polamuri .
4. Cross Entropy
 - a. **“A Short Introduction to Entropy, Cross-Entropy and KL-Divergence”**, Aurélien Géron. <https://www.youtube.com/watch?v=ErfnhcEV1O8> .
 - b. **A Friendly Introduction to Cross-Entropy Loss**, Rob DiPietro. <https://rdipietro.github.io/friendly-intro-to-cross-entropy-loss/>.